

Nominatim Apptainer Container

- Nominatim SQLite Geocoding

Version: Nominatim 5.1.0

Base: Ubuntu 24.04

Output: Portable SQLite Database

Summary:

A self-contained geocoding service that:

- Imports OpenStreetMap data using PostgreSQL
- Converts to SQLite for portability
- Includes Iowa OSM data + US metadata
- Runs without external dependencies

Use Case: Address lookup, reverse geocoding, location search

Why PostgreSQL During Build?

No Alternative (yet): Nominatim's import pipeline is PostgreSQL-only

Convert to SQLite?

Benefits:

- Single File
- Zero Dependencies: No PostgreSQL daemon needed
- May be easier to have multiple nodes batch on Argon.
- Avoids having to deal with multiple namespaces post build
 - pguser, nominatim, and www-data

Trade-off:

- No concurrent writes (but our geocoding is read-only!)

/app vs /nominatim:

During the build we are mainly operating inside of two directories:

/app/

/app - Build-time workspace

- Temporary files needed only during container build
- PostgreSQL database and runtime files
- Configuration files
- Downloaded OSM and metadata files
- Removed in cleanup phase to reduce final image size

/nominatim/

/nominatim - Runtime workspace

- Nominatim project directory
- Final SQLite database (nominatim.sqlite)
- .env configuration file
- Kept in final container for geocoding queries

Container File Structure

```
nominatim-container/
    └── nominatim.def                                # Main container definition file
    └── conf.d/
        ├── postgresql.conf                          # PostgreSQL & Nominatim configuration
        ├── pg_hba.conf                             # PostgreSQL performance settings
        ├── postgres-import.conf                    # Database authentication rules
        ├── postgres-tuning.conf                   # Import-time optimizations
        └── env                                    # Runtime performance tuning
            # Nominatim environment variables
    └── nominatim_project/
        ├── config.sh                            # Data files and scripts
        ├── iowa-latest.osm.pbf                  # Nominatim configuration script
        ├── secondary_importance.sql.gz          # OpenStreetMap data (50-200MB)
        ├── us_postcodes.csv.gz                 # Additional search ranking
        ├── wikimedia-importance.csv.gz         # US ZIP code database
        └── tiger-nominatim-preprocessed-latest.csv.tar.gz # Search result weights
                                                # US Census addresses
```

Required Files

Configuration (`conf.d/`):

- `postgresql.conf` - Performance tuning
- `pg_hba.conf` - Authentication
- `postgres-import.conf` - Import optimizations
- `env` - Environment variables

Required Files

Data (`nominatim_project/`):

- `iowa-latest.osm.pbf` - OpenStreetMap data
- `us_postcodes.csv.gz` - ZIP codes
- `wikimedia-importance.csv.gz` - Search ranking
- `tiger-nominatim-preprocessed-latest.csv.tar.gz` - US addresses

Build Process: Phase 1

System Setup

APT Packages:

```
apt-get install \
  postgresql-postgis \
  osm2pgsql \
  python3-pip \
  gdal-bin
```

Python Packages:

```
pip install --break-system-packages \
  nominatim-db==5.1.0 \
  nominatim-api==5.1.0
```

Phase 2: User & Database Setup

- Create non-root PostgreSQL User
- Initialize Database

```
useradd -m -s /bin/bash pguser  
su pguser -c "initdb -D $PGDATA --auth=trust"
```

Why **--auth=trust**?

- Ran into errors during build with users
- Only relevant during build time so no real network exposure when doing geocoding.

Database Users:

Nominatim requires the following users:

```
createuser --superuser pguser          # Owns PostgreSQL database
createuser --no-superuser nominatim     # Read-only
createuser --no-superuser www-data       # Web accessor
```

Phase 3: Nominatim Import

```
nominatim import \
--osm-file /app/data/iowa-latest.osm.pbf \
--project-dir /nominatim \
--threads 4
```

- People say to use ~1 thread per 2GB of RAM. For my computer this is fine.

There is a `.env` file which Nominatim uses for build variables.

This along with postgresql conf will probably need to be optimized.

Phase 4: Metadata

Import to Nominatim

- US Postcodes (ZIP code geocoding)
- Wikimedia Importance (improves result ranking)
- Tiger-Line Data (address interpolation)

Phase 5: SQLite Convert

Experimental

```
nominatim convert -o /nominatim/nominatim.sqlite
```

Converts placex table, search indexes, address hierarchies, and metadata tables

Phase 6: Cleanup

- Stop PostgreSQL
- Remove Build Artifacts

Mainly done to reduce the size of the resulting container to just have SQLite database.

Offline Container Build Strategy:

Problem:

Building containers offline (Argon) requires all dependencies be pre-downloaded and no elevated privileges.

- APT Packages
- Python wheels
- Base Docker images

Solution: Two Stage Build

To run privileged workflows on Argon we need an apptainer container.

So to build the nominatim.sif we will first create a dependency container which will be used as a cache when building the actual container.

1. `resources.def` container which packages all dependencies
2. `builder.def` Container which uses resources as a base to create a nominatim.sif

Overview:

Stage 1: resources.def (Build requires internet)

Bootstrap: docker
From: ubuntu:24.04

- Downloads & installs all APT packages
- Downloads & installs all Python wheels
- Caches dependencies for offline use

Output: resources.sif

and

Stage 2: builder.def (offline)

Bootstrap: localimage
From: resources.sif

- Uses pre-installed dependencies
- Imports OSM data
- Builds PostgreSQL database
- Converts to SQLite

Output: nominatim.sif

builder.def Demo

```
Bootstrap: docker
From: ubuntu:24.04
```

```
%files
offline_deps/apt_packages/*.deb /opt/apt_packages/
offline_deps/python_wheels/*.whl /opt/python_wheels/

%post
# Install all dependencies
dpkg -i /opt/apt_packages/*.deb
pip install --no-index --find-links=/opt/python_wheels \
nominatim-db nominatim-api osmium psycopg aiosqlite
```

application.def Demo

```
Bootstrap: localimage      # Uses local .sif file
From: builder.sif

%files
    nominatim_project/iowa-latest.osm.pbf /app/data/

%post
    # All dependencies already installed
    nominatim import --osm-file /app/data/iowa-latest.osm.pbf
    nominatim convert -o /nominatim/nominatim.sqlite
```

Will be very similar to nominatim.def but instead pulls dependencies from builder.def.